

Euclidean Heuristic Optimization

Chris Rayner (University of Alberta)
Michael Bowling (University of Alberta)
Nathan Sturtevant (University of Denver)

AAAI 2011

August 9, 2011



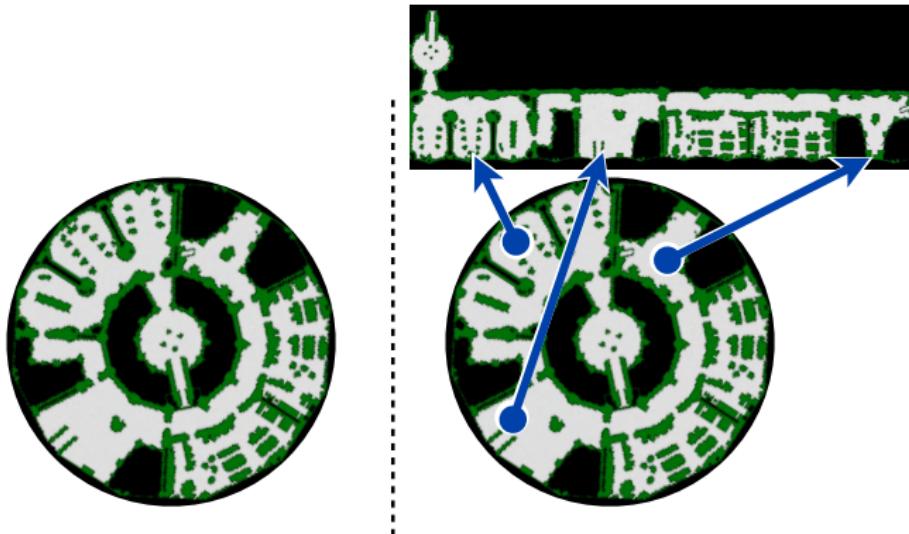
Build memory-efficient data structures *before shipping* to help:

- ▶ Solve “easy” problems in rapid succession
- ▶ Any state can be the goal



Dragon Age: Origins (BioWare)

- ▶ Straight-line heuristics are easy to compute *online*



Dragon Age: Origins (BioWare)

- ▶ Straight-line heuristics are easy to compute *online*
- ▶ Improve them by (re-)arranging the search graph *offline*

(Example 1) Rearranging a grid world

(Example 2) Arranging a 2×3 sliding tile puzzle:



EUCLIDEAN HEURISTICS

Definition (Euclidean heuristics)

A Euclidean heuristic Y is one whose heuristic values can be computed as distances in a Euclidean space of d dimensions

$$h(i, j) = \|y_i - y_j\|$$

EUCLIDEAN HEURISTICS

Definition (Euclidean heuristics)

A Euclidean heuristic Y is one whose heuristic values can be computed as distances in a Euclidean space of d dimensions

$$h(\textcolor{red}{i}, j) = \|\textcolor{red}{y_i} - y_j\|$$

EUCLIDEAN HEURISTICS

Definition (Euclidean heuristics)

A Euclidean heuristic Y is one whose heuristic values can be computed as distances in a Euclidean space of d dimensions

$$h(\textcolor{red}{i}, \textcolor{blue}{j}) = \|y_i - y_j\|$$

PROBLEM STATEMENT

Look among *all* Euclidean heuristics Y for one that is best

PROBLEM STATEMENT

Look among *all* Euclidean heuristics Y for one that is best

Definition (Optimal Euclidean Heuristic)

Minimizes the loss between true distances δ and heuristics h :

$$\underset{Y}{\text{minimize}} \quad \mathcal{L}(Y)$$

subject to Y is admissible and consistent

PROBLEM STATEMENT

Look among *all* Euclidean heuristics Y for one that is best

Definition (Optimal Euclidean Heuristic)

Minimizes the loss between true distances δ and heuristics h :

$$\underset{Y}{\text{minimize}} \quad \mathcal{L}(Y)$$

subject to Y is **admissible and consistent**

ADMISSIBILITY / CONSISTENCY

SIMPLIFIED CONSTRAINTS

A Euclidean heuristic Y is **admissible** and **consistent** if:

$$\forall(i, j) \quad \|y_i - y_j\| \leq \delta(i, j)$$

$$\forall(i, j, k) \quad \|y_i - y_j\| \leq \delta(i, k) + \|y_j - y_k\|$$

ADMISSIBILITY / CONSISTENCY

SIMPLIFIED CONSTRAINTS

A Euclidean heuristic Y is **admissible** and **consistent** if:

$$\forall(i, j) \quad \|y_i - y_j\| \leq \delta(i, j)$$

$$\forall(i, j, k) \quad \|y_i - y_j\| \leq \delta(i, k) + \|y_k - y_j\|$$

Theorem

Y is locally admissible

$$\forall(i, j) \in E \quad \|y_i - y_j\| \leq \delta(i, j)$$

\Updownarrow

Y is admissible and consistent

Passino and Antsaklis, 1994

ADMISSIBILITY / CONSISTENCY

SIMPLIFIED CONSTRAINTS

A Euclidean heuristic Y is **admissible** and **consistent** if:

$$\forall(i, j) \quad \|y_i - y_j\| \leq \delta(i, j)$$

$$\forall(i, j, k) \quad \|y_i - y_j\| \leq \delta(i, k) + \|y_k - y_j\|$$

Theorem

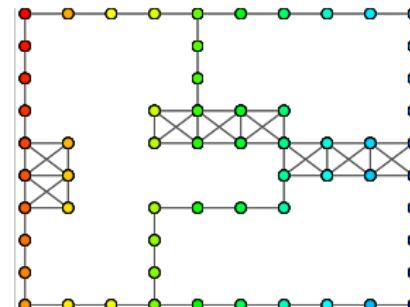
Y is locally admissible

$$\forall(i, j) \in E \quad \|y_i - y_j\| \leq \delta(i, j)$$



Y is admissible and consistent

Passino and Antsaklis, 1994



(one constraint per edge)

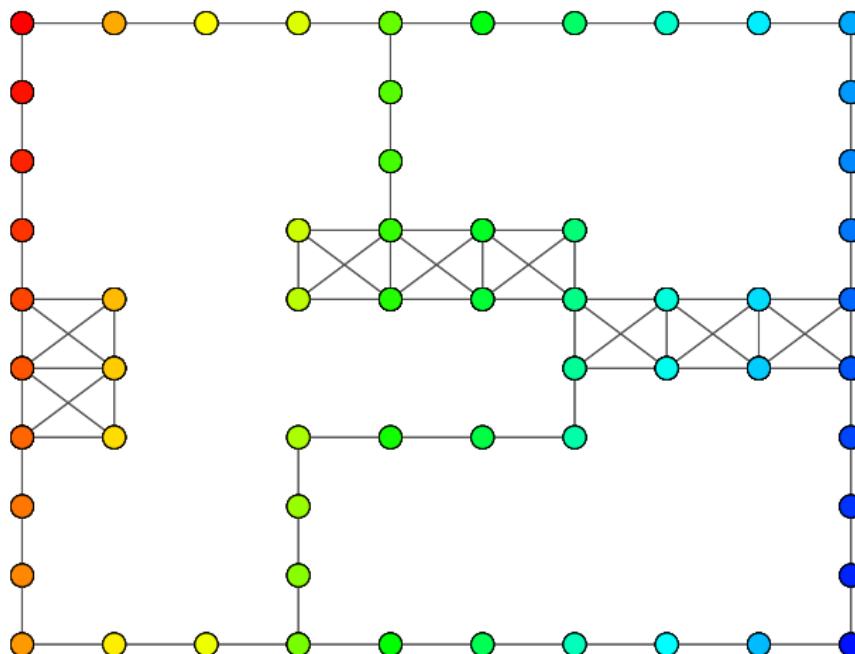
Definition (Optimal Euclidean Heuristic)

Minimizes the **loss** between true distances δ and heuristics h :

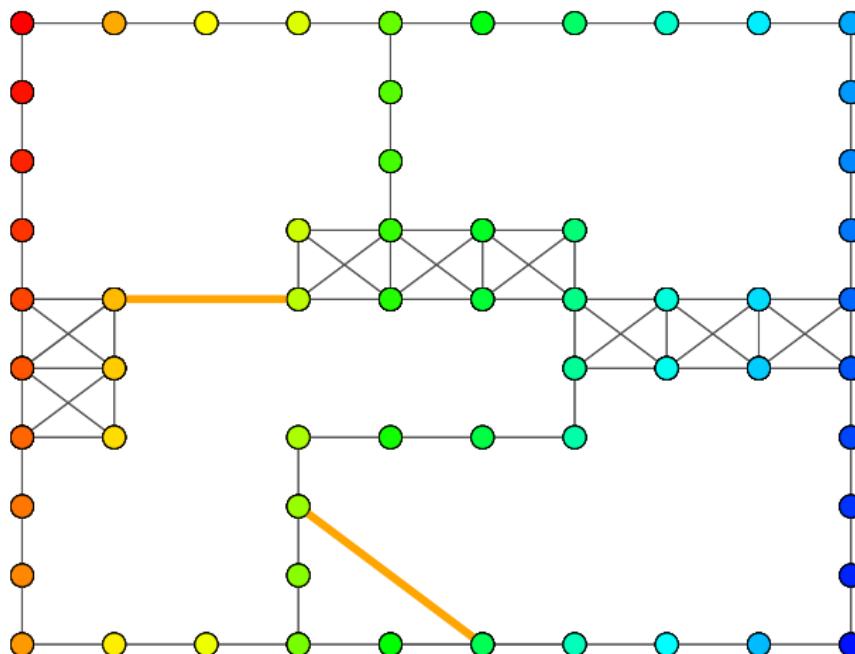
$$\underset{Y}{\text{minimize}} \quad \mathcal{L}(Y)$$

subject to Y is admissible and consistent

DEFINING LOSS

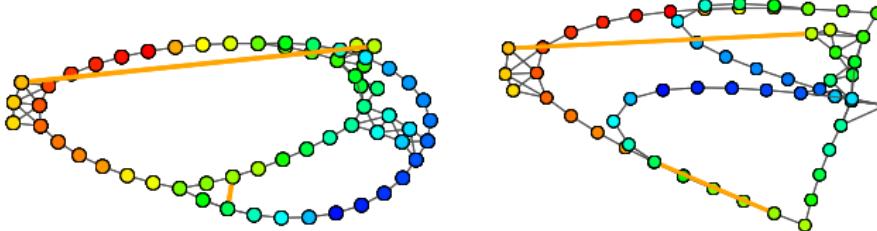


DEFINING LOSS



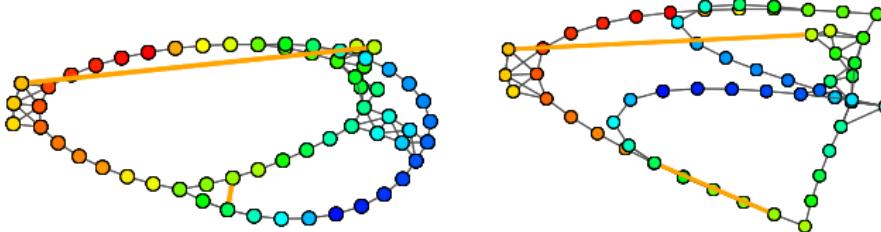
DEFINING LOSS

Which is better?



DEFINING LOSS

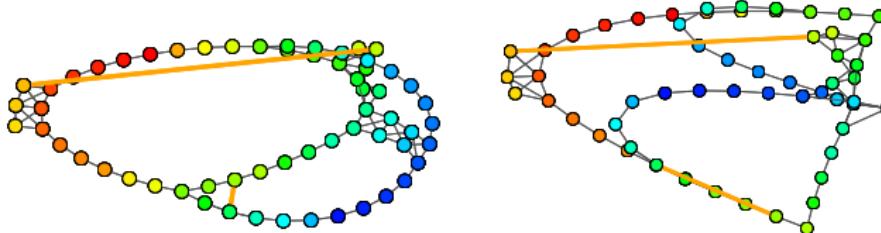
Which is better?



Loss function \mathcal{L} combines errors $\forall(i,j)$ into a single scalar

DEFINING LOSS

Which is better?

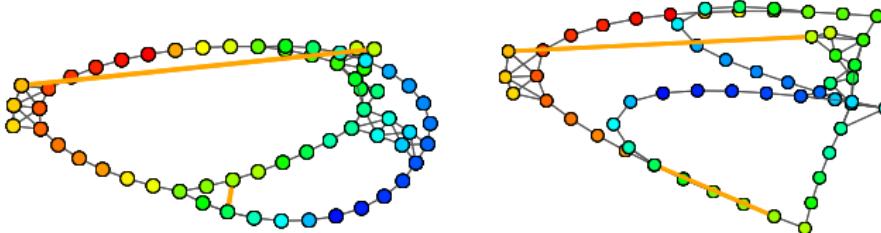


Loss function \mathcal{L} combines errors $\forall(i, j)$ into a single scalar

- ▶ specify trade-off: many small errors *vs.* a large error?

DEFINING LOSS

Which is better?

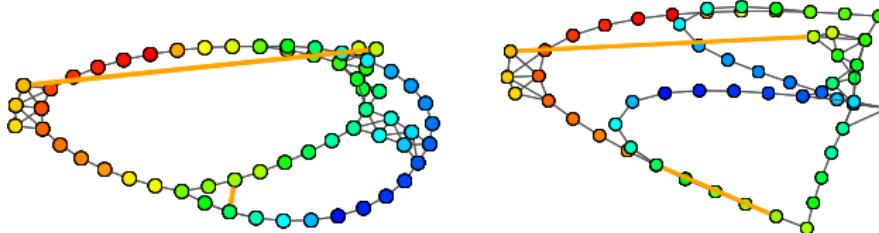


Loss function \mathcal{L} combines errors $\forall(i, j)$ into a single scalar

- ▶ specify trade-off: many small errors *vs.* a large error?
- ▶ specify relative importance of each state pair

DEFINING LOSS

Which is better?



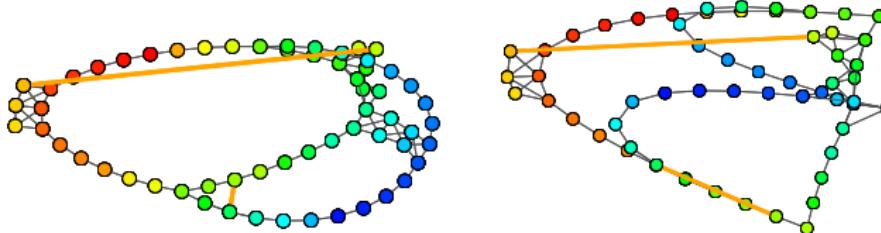
Loss function \mathcal{L} combines errors $\forall(i,j)$ into a single scalar

- ▶ specify trade-off: many small errors *vs.* a large error?
- ▶ specify relative importance of each state pair

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

DEFINING LOSS

Which is better?



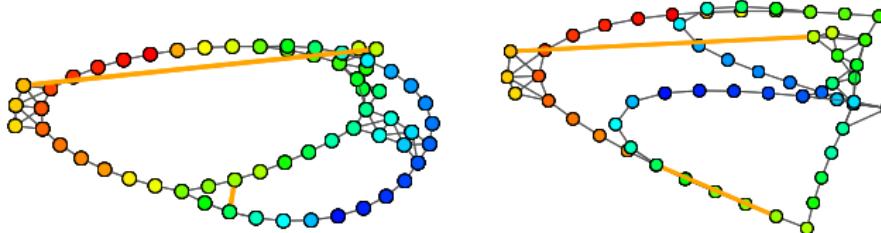
Loss function \mathcal{L} **combines** errors $\forall(i,j)$ into a single scalar

- ▶ specify trade-off: many small errors *vs.* a large error?
- ▶ specify relative importance of each state pair

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

DEFINING LOSS

Which is better?



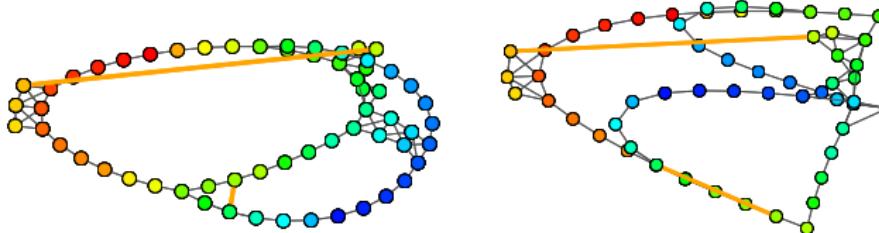
Loss function \mathcal{L} combines errors $\forall(i,j)$ into a single scalar

- ▶ **specify trade-off:** many small errors *vs.* a large error?
- ▶ specify relative importance of each state pair

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

DEFINING LOSS

Which is better?



Loss function \mathcal{L} combines errors $\forall(i,j)$ into a single scalar

- ▶ specify trade-off: many small errors *vs.* a large error?
- ▶ **specify relative importance** of each state pair

$$\mathcal{L}(Y) = \sum_{i,j} \textcolor{red}{W_{ij}} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

$$\text{Admissibility} \rightarrow \delta(i,j)^2 \geq \|y_i - y_j\|^2$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

Admissibility $\rightarrow \delta(i,j)^2 \geq \|y_i - y_j\|^2$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} (\delta(i,j)^2 - \|y_i - y_j\|^2)$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

Admissibility $\rightarrow \delta(i,j)^2 \geq \|y_i - y_j\|^2$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} (\delta(i,j)^2 - \|y_i - y_j\|^2)$$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} \delta(i,j)^2 - \sum_{i,j} W_{ij} \|y_i - y_j\|^2$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

Admissibility $\rightarrow \delta(i,j)^2 \geq \|y_i - y_j\|^2$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} (\delta(i,j)^2 - \|y_i - y_j\|^2)$$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} \delta(i,j)^2 - \sum_{i,j} W_{ij} \|y_i - y_j\|^2$$

DEFINING LOSS

SIMPLIFYING THE OBJECTIVE

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} |\delta(i,j)^2 - \|y_i - y_j\|^2|$$

Admissibility $\rightarrow \delta(i,j)^2 \geq \|y_i - y_j\|^2$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} (\delta(i,j)^2 - \|y_i - y_j\|^2)$$

$$\mathcal{L}(Y) = \sum_{i,j} W_{ij} \delta(i,j)^2 - \sum_{i,j} W_{ij} \|y_i - y_j\|^2$$

$$\boxed{\mathbf{maximize}_Y \sum_{i,j} W_{ij} \|y_i - y_j\|^2}$$

CONNECTIONS

NONLINEAR DIMENSIONALITY REDUCTION

The full optimization problem:

$$\begin{aligned} & \underset{Y}{\text{maximize}} && \sum_{i,j} W_{ij} \|y_i - y_j\|^2 \\ & \text{subject to} && \forall (i, j) \in E \quad \|y_i - y_j\| \leq d(i, j) \end{aligned}$$

CONNECTIONS

NONLINEAR DIMENSIONALITY REDUCTION

The full optimization problem:

$$\begin{aligned} & \underset{Y}{\text{maximize}} && \sum_{i,j} W_{ij} \|y_i - y_j\|^2 \\ & \text{subject to} && \forall (i, j) \in E \quad \|y_i - y_j\| \leq d(i, j) \end{aligned}$$

This is a weighted generalization of Weinberger *et al.*'s
Maximum Variance Unfolding (MVU)

CONNECTIONS

NONLINEAR DIMENSIONALITY REDUCTION

The full optimization problem:

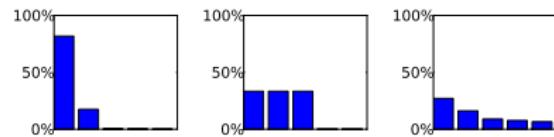
$$\begin{aligned} & \underset{Y}{\text{maximize}} && \sum_{i,j} W_{ij} \|y_i - y_j\|^2 \\ & \text{subject to} && \forall (i, j) \in E \quad \|y_i - y_j\| \leq d(i, j) \end{aligned}$$

This is a weighted generalization of Weinberger *et al.*'s
Maximum Variance Unfolding (MVU)

Heuristic learning is linked to manifold learning

MANIFOLD LEARNING

- ▶ What “dimensionality” will hold a search graph?
- ▶ Visualization may shed new light on problems



DIFFERENTIAL HEURISTICS

NG & ZHANG '02, GOLDBERG & HARRELSON '05, STURTEVANT *et al.*'09

Imagine “hooking” the graph on a **pivot** state

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

Let “pivot” be state 1 and $n = 4$:

$$W_{\text{diff}} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & -1/3 & -1/3 \\ 1 & -1/3 & 0 & -1/3 \\ 1 & -1/3 & -1/3 & 0 \end{bmatrix}$$

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

Let “pivot” be state 1 and $n = 4$:

$$W_{\text{diff}} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & -1/3 & -1/3 \\ 1 & -1/3 & 0 & -1/3 \\ 1 & -1/3 & -1/3 & 0 \end{bmatrix}$$

- ▶ Push points *away* from the pivot state (**weight 1**)

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

Let “pivot” be state 1 and $n = 4$:

$$W_{\text{diff}} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & -1/3 & -1/3 \\ 1 & -1/3 & 0 & -1/3 \\ 1 & -1/3 & -1/3 & 0 \end{bmatrix}$$

- ▶ Push points *away* from the pivot state (weight 1)
- ▶ and pull points *into* each other (weight $1/(n-1)$)

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

Let “pivot” be state 1 and $n = 4$: Can this be improved?

$$W_{\text{diff}} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & -1/3 & -1/3 \\ 1 & -1/3 & 0 & -1/3 \\ 1 & -1/3 & -1/3 & 0 \end{bmatrix}$$

- ▶ Push points *away* from the pivot state (weight 1)
- ▶ and pull points *into* each other (weight $1/(n-1)$)

CONNECTIONS

DIFFERENTIAL HEURISTICS



Our approach recovers differential heuristics when $W = W_{\text{diff}}$:

Let “pivot” be state 1 and $n = 4$: Can this be improved?

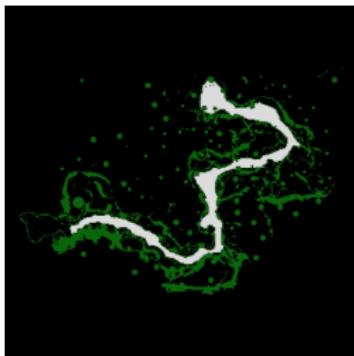
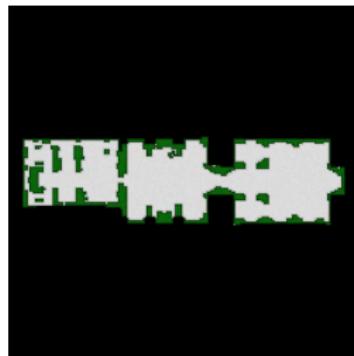
$$W_{\text{diff}} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & -1/3 & -1/3 \\ 1 & -1/3 & 0 & -1/3 \\ 1 & -1/3 & -1/3 & 0 \end{bmatrix}$$

$$W_{\text{diff}}^+ = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & \epsilon & \epsilon \\ 1 & \epsilon & 0 & \epsilon \\ 1 & \epsilon & \epsilon & 0 \end{bmatrix}$$

- ▶ Push points *away* from the pivot state (weight 1)
- ▶ and pull points *into* each other (weight $1/(n-1)$)

Let's experiment with $\epsilon = 10^{-3}$...

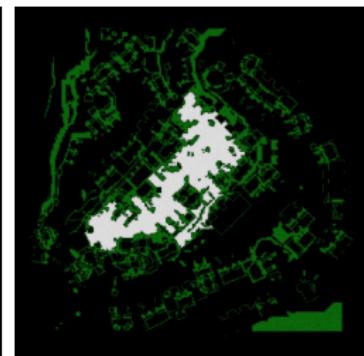
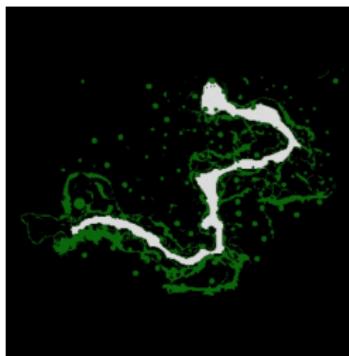
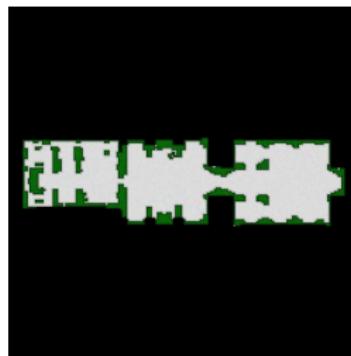
PATHPLANNING EXPERIMENT



Dragon Age: Origins (BioWare)
www.movingai.com

- ▶ Maps with 168–6,240 states: standard problem sets

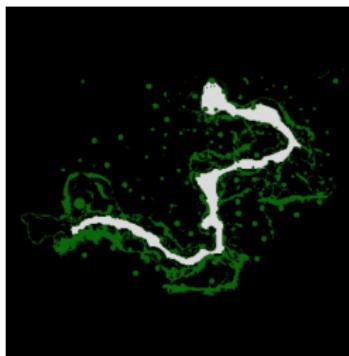
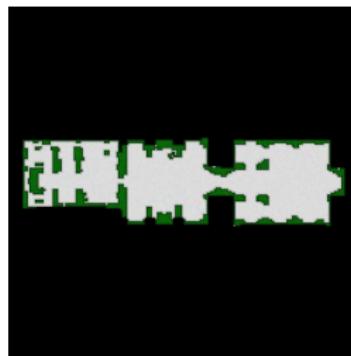
PATHPLANNING EXPERIMENT



Dragon Age: Origins (BioWare)
www.movingai.com

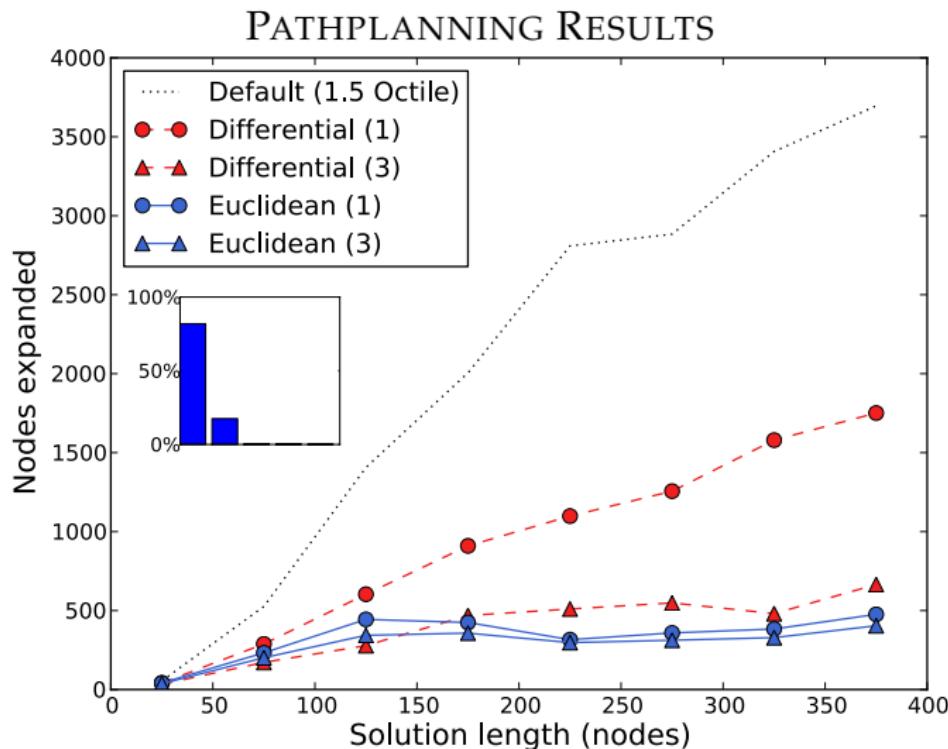
- ▶ Maps with 168–6,240 states: standard problem sets
- ▶ Count the nodes A* [Hart *et al.* 68] expands to find a path

PATHPLANNING EXPERIMENT



Dragon Age: Origins (BioWare)
www.movingai.com

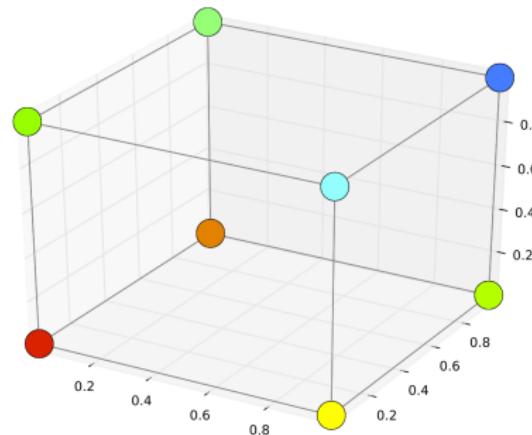
- ▶ Maps with 168–6,240 states: standard problem sets
- ▶ Count the nodes A* [Hart *et al.* 68] expands to find a path
- ▶ Compare W_{diff} and W_{diff}^+ :



Sets of 1 and 3 differential heuristics *vs.*
Sets of 1 and 3 Euclidean heuristics ($W = W_{\text{diff}}$)

CUBE WORLD EXPERIMENT

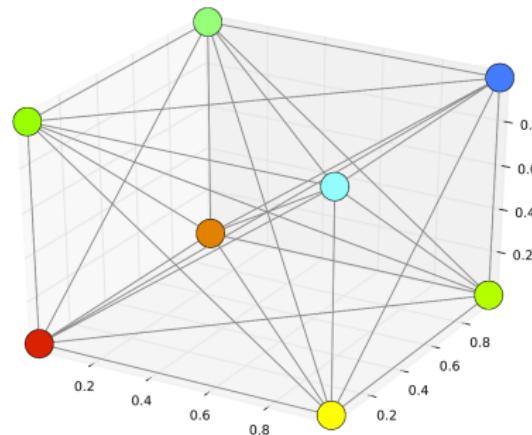
Differential heuristics' weakness: high dimensionality



- ▶ Octile grid world, generalized to higher dimensions

CUBE WORLD EXPERIMENT

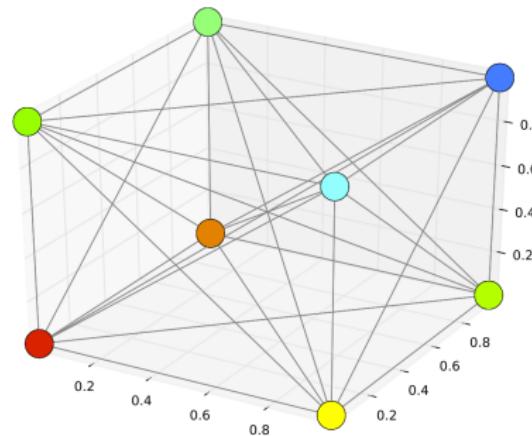
Differential heuristics' weakness: high dimensionality



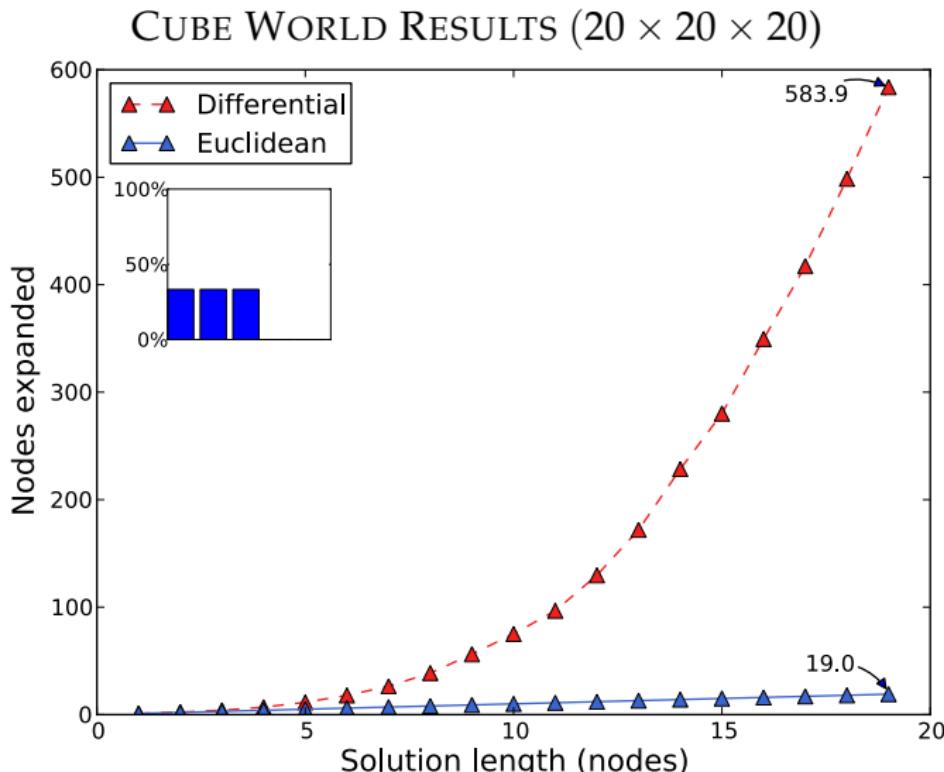
- ▶ Octile grid world, generalized to higher dimensions
 - ▶ Agent can increment any/all coordinates by 1 each turn

CUBE WORLD EXPERIMENT

Differential heuristics' weakness: high dimensionality



- ▶ Octile grid world, generalized to higher dimensions
 - ▶ Agent can increment any/all coordinates by 1 each turn
- ▶ Transition costs are edge lengths



Set of 4 differential heuristics *vs.*
one 3-dimensional Euclidean heuristic ($W_{ij} = 1$)

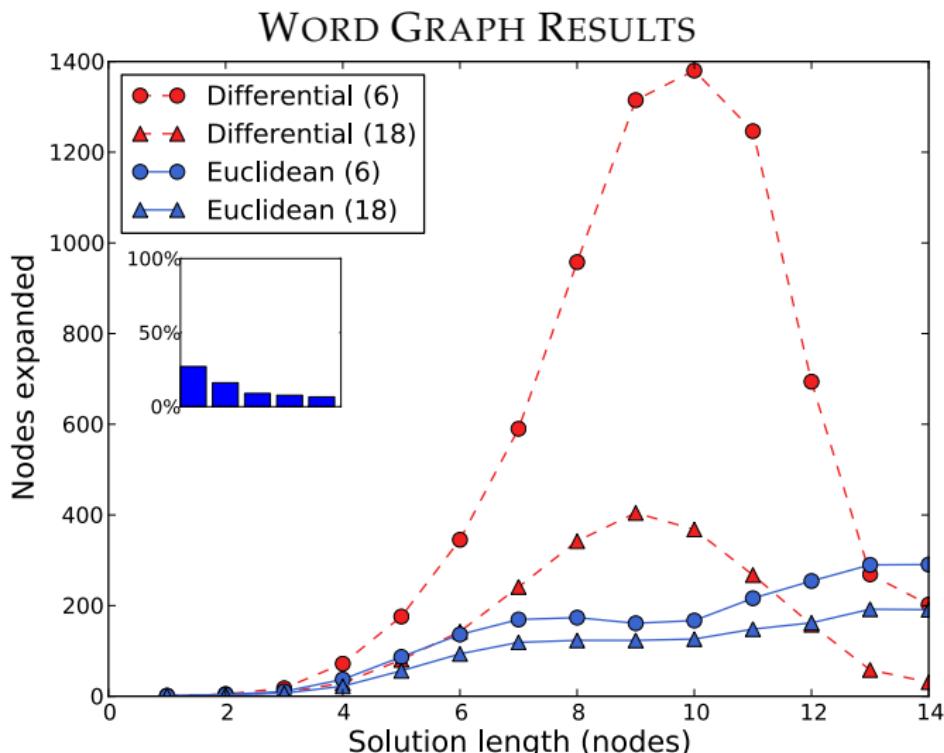
WORD GRAPH EXPERIMENT

- ▶ 4,820 states representing four-letter words

WORD GRAPH EXPERIMENT

- ▶ 4,820 states representing four-letter words
- ▶ Find shortest sequence of 1-letter changes turning start word into goal word

fore → fork → ? → back



Sets of 6 and 18 differential heuristics *vs.*
6 and 18-dimensional Euclidean heuristics ($W_{ij} = 1$)

SUMMARY

Euclidean Heuristic Optimization

A novel way to build admissible/consistent heuristics

- ▶ principled link to manifold learning
- ▶ generalization of differential heuristics
- ▶ promising empirical results on small problems

(Thank you: Ariel Felner, our anonymous reviewers,
NSERC and iCore)